

OpenRAG Knowledge Base for Automated Client Content Generation

A RAG-powered assistant that turns IBM's internal content library into instant, grounded, client-ready sales assets — built on IBM Bob, OpenRAG, and watsonx.ai.

Team: 4Bs (3Boys + 1Bob) **Track:** Sales Enablement / Generative AI **Stack:** IBM Bob · OpenRAG · watsonx.ai · Box SDK · Seismic

THE PROBLEM

Every deal wastes hours on content that already exists

IBM's approved docs, case studies, and battlecards exist — scattered across Box, Seismic, and email. Every time a consultant needs a one-pager, they start from scratch: searching multiple platforms, verifying versions, manually stitching content. That costs **2 to 4 hours per asset, per deal**.

The result: outdated content reaches clients, institutional knowledge is locked in individuals, and consultants burn time on assembly instead of selling.

TARGET USER

Client-facing roles who need the right content, right now

Stakeholder	Need	Pain
Sales Consultant (Primary)	One-pagers & deck outlines per deal	2–4 h manual assembly per asset
Sales Engineer	Technical briefs on demand	No fast retrieval mid-call
Client Partner	Current, on-brand collateral	Relies on individual memory

THE SOLUTION

Type a question. Get a client-ready asset in 20 seconds.

We are building an OpenRAG knowledge base that continuously ingests approved IBM content from Box and Seismic, indexes it semantically, and uses IBM Bob to generate polished, source-cited client assets on demand. No searching. No copy-pasting. No stale documents.

Ingestion: Box + Seismic → Fetch & Parse → Chunk → OpenRAG (embed + index)

Generation: User Query → IBM Bob (rewrite) → OpenRAG (retrieve) → IBM Bob / Granite (generate) → Client Asset with citations

Every output is grounded in approved source content with inline citations. If the answer is not in the knowledge base, the system says so — no hallucinations, no guesswork. In practice: "Generate a one-pager for a retail client interested in watsonx.ai." Done in 20 seconds. Hours of work handed back to the team, every single day.

HOW WE WILL USE IBM BOB

Bob is not just a tool we use — it is what makes this system work

- **Query rewriting:** Turns vague requests into precise retrieval queries so OpenRAG returns the most relevant content every time.
- **Asset generation:** Synthesises retrieved chunks into finished one-pagers, deck outlines, and summaries using prompt templates tuned for each output type.
- **Prompt engineering:** Bob's built-in skill system lets us iterate on generation quality directly in the dev environment — fast feedback, no context switching.
- **Competitive Q&A:** Answers open-ended positioning questions grounded in battlecards retrieved live from the knowledge base.
- **Code scaffolding:** Generates ingestion pipeline boilerplate — Box SDK, Seismic integration, OpenRAG collection management — cutting build time significantly.
- **Security:** Credential guidance and IBM Key Protect integration ensure no API keys are ever hardcoded anywhere in the project.

EXPECTED IMPACT

3 h → 20 s Time to first-draft client asset	100% Outputs grounded in approved IBM content	10+ Active users within 4 weeks of launch
---	---	---